

Contribution Statement

All four members contributed equally (25% each). Everyone was involved in studying the baseline papers, designing the algorithm, running experiments, and writing the report.

Vedant Jawandhia (2022B4A70627P): Wrote the core PURGE code — gradient projection loop, kl_retain objective, KD anchoring, and the main unlearning pipeline. Handled hyperparameter tuning and ran the CIFAR-10 and cross-dataset experiments. Wrote the methodology section of the report.

Chavi Gupta (2022B3A70637P): Co-implemented the unlearning pipeline — worked on representation erasure, BatchNorm freezing logic, and the evaluation code (MIA AUROC, feature-level metrics). Ran the PathMNIST and MNIST experiments, tracked ablation results, and generated the result tables and figures. Wrote parts of the results and ablation sections.

Daksh Ahuja (2022A7PS0123P): Studied A-GEM, SalUn, and gradient ascent baselines. Wrote the related work and limitations sections. Put together the claim-framing documents and helped frame the CL-MU duality narrative. Prepared viva talking points.

Samyak Savi (2022B3A70635P): Studied KD-based unlearning, Fisher-based methods, and MIA evaluation methodology. Set up the repo structure, wrote shell scripts and YAML configs, handled submission packaging. Did code review, reproducibility checks, and co-created the demo video.

All members reviewed the final report and verified the claims against code and logs.